✖

Ridge Regression

9 试题

1.

Which of the following is NOT a valid measure of overfitting?

○ <mark>Sum of parameters (w1+w2+...+wn)</mark>

○ Sum of squares of parameters (w1^2 + w2^2 + ... +wn^2)

○ Range of parameters, i.e., difference between maximum and minimum parameters

○ Sum of absolute values of parameters (|w1| + |w2| + ... + |wn|)
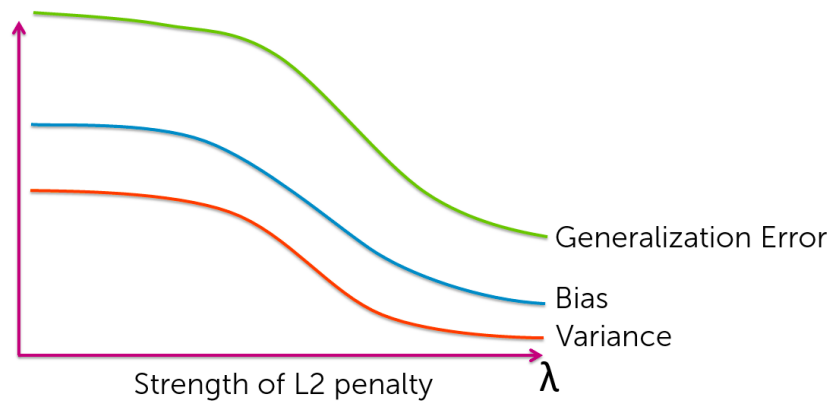

2.

In ridge regression, choosing a large penalty strength λ tends to lead to a model with (choose all that apply):

☐ <mark>High bias</mark>
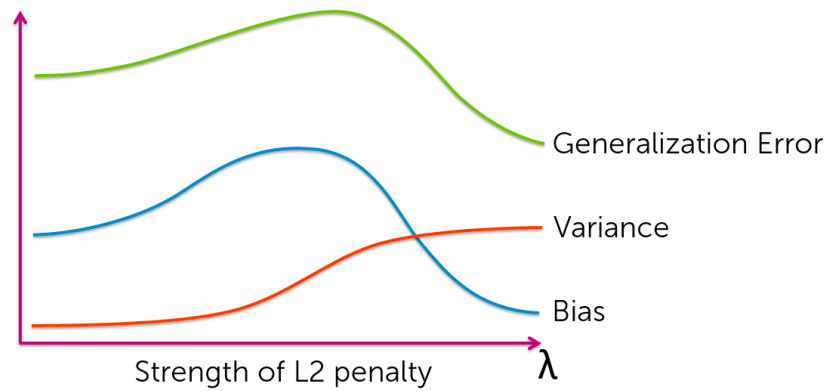
☐ Low bias

☐ High variance

☐ <mark>Low variance</mark>

3.

Which of the following plots best characterize the trend of bias, variance, and generalization error (all plotted over λ)?
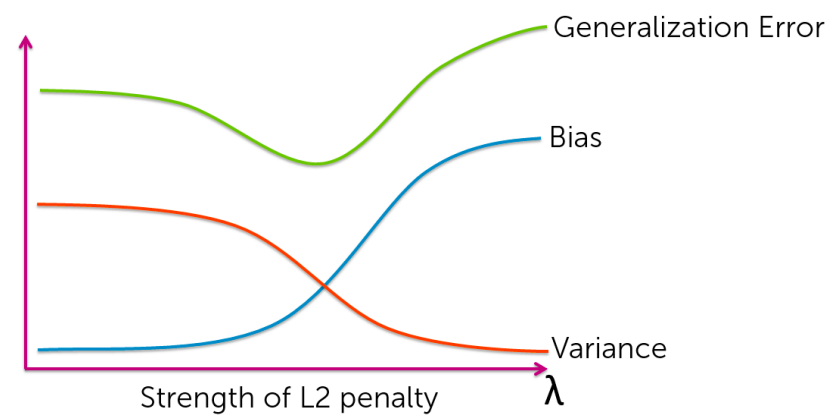
○

4.

○

💬

4.

In ridge regression using unnormalized features, if you double the value of a given feature (i.e., a specific column of the feature matrix), what happens to the estimated coefficients for every other feature? They:

○ Double

○ Half

○ Stay the same

○ ==Impossible to tell from the information provided==

5.

If we only have a small number of observations, K-fold cross validation provides a better estimate of the generalization error than the validation set method.

○ ==True==

○ False

6.

10-fold cross validation is more computationally intensive than leave-one-out (LOO) cross validation.

○ True

○ ==False==

7.

Assume you have a training dataset consisting of N observations and D features. You use the closed-form solution to fit a multiple linear regression model using ridge regression. To choose the penalty strength λ, you run leave-one-out (LOO) cross validation searching over L values of λ. Let Cost(N,D) be the computational cost of running ridge regression with N data points and D features. Assume the prediction cost is negligible compared to the computational cost of training the model. Which of the following represents the computational cost of your LOO cross validation procedure?

○   L* N * Cost(N,D)

○   <mark>L* N * Cost(N-1,D)</mark>

○   L* D * Cost(N-1,D)

○   L* D * Cost(N,D)

○   L* Cost(N-1,D)

○   L* Cost(N,D)

8.
Assume you have a training dataset consisting of 1 million observations. Suppose running the closed-form solution to fit a multiple linear regression model using ridge regression on this data takes 1 second. Suppose you want to choose the penalty strength λ by searching over 100 possible values. How long will it take to run leave-one-out (LOO) cross-validation for this selection task?

○   About 3 hours

○   About 3 days

○   <mark>About 3 years</mark>

○   About 3 decades

9.
Assume you have a training dataset consisting of 1 million observations. Suppose running the closed-form solution to fit a multiple linear regression model using ridge regression on this data takes 1 second. Suppose you want to choose the penalty strength λ by searching over 100 possible values. If you only want to spend about 1 hour to select λ, what value of k should you use for k-fold cross-validation?

○   k=6

○   <mark>k=36</mark>

○   k=600

○   k=3600

提交测试

👍   👎   🏳